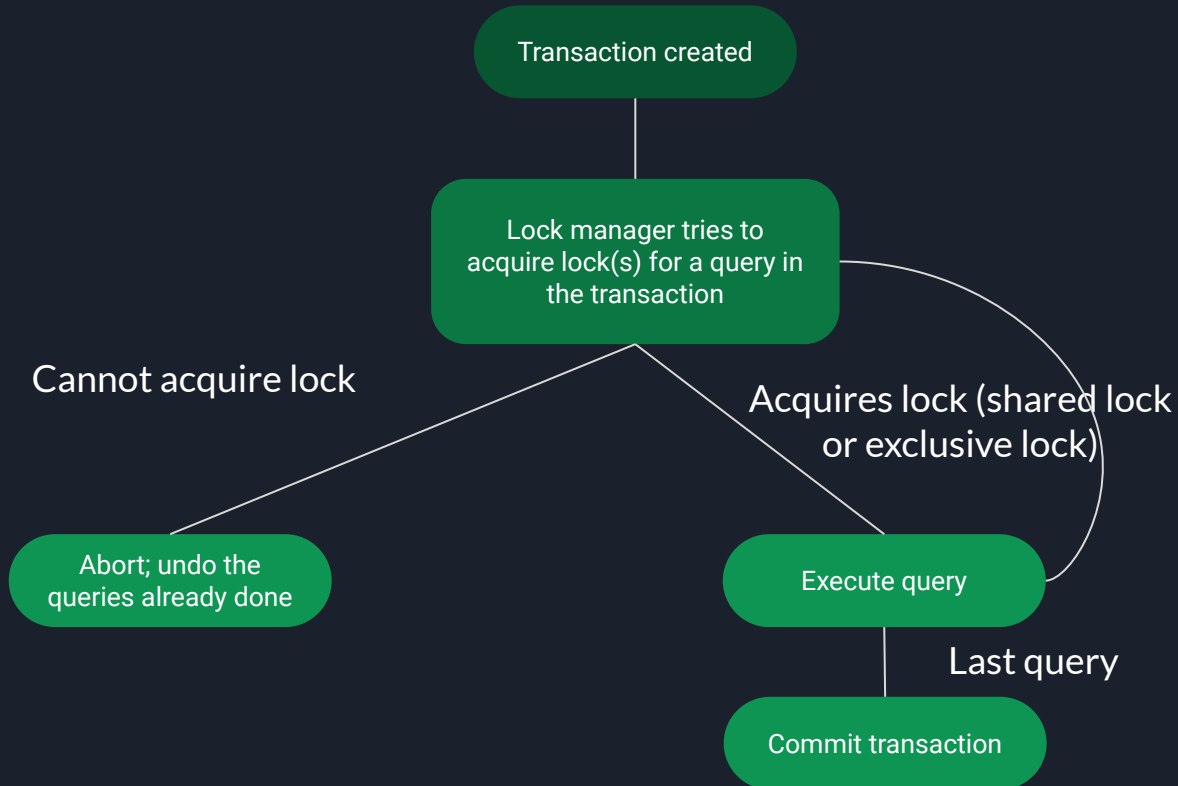




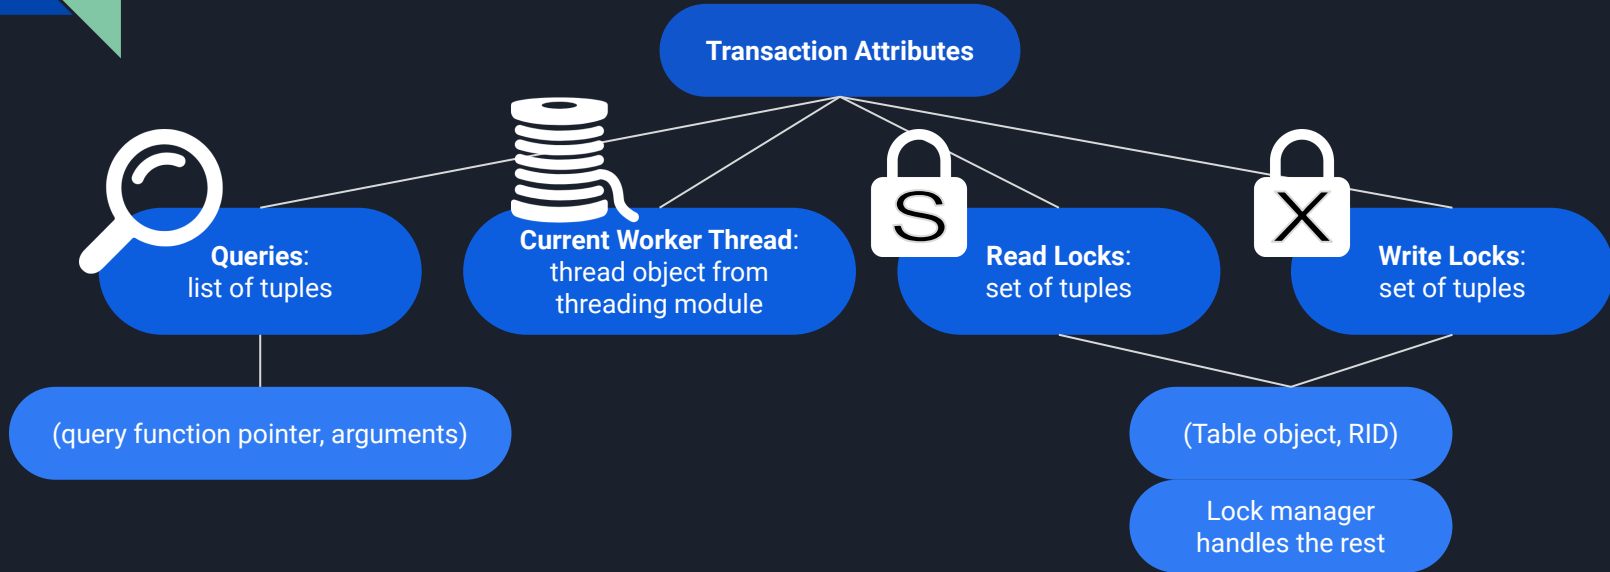
ECS 165A Milestone 3

Aditi Bali, Jie Lyu, Jessica Ma, Nathan Chan, William Shih

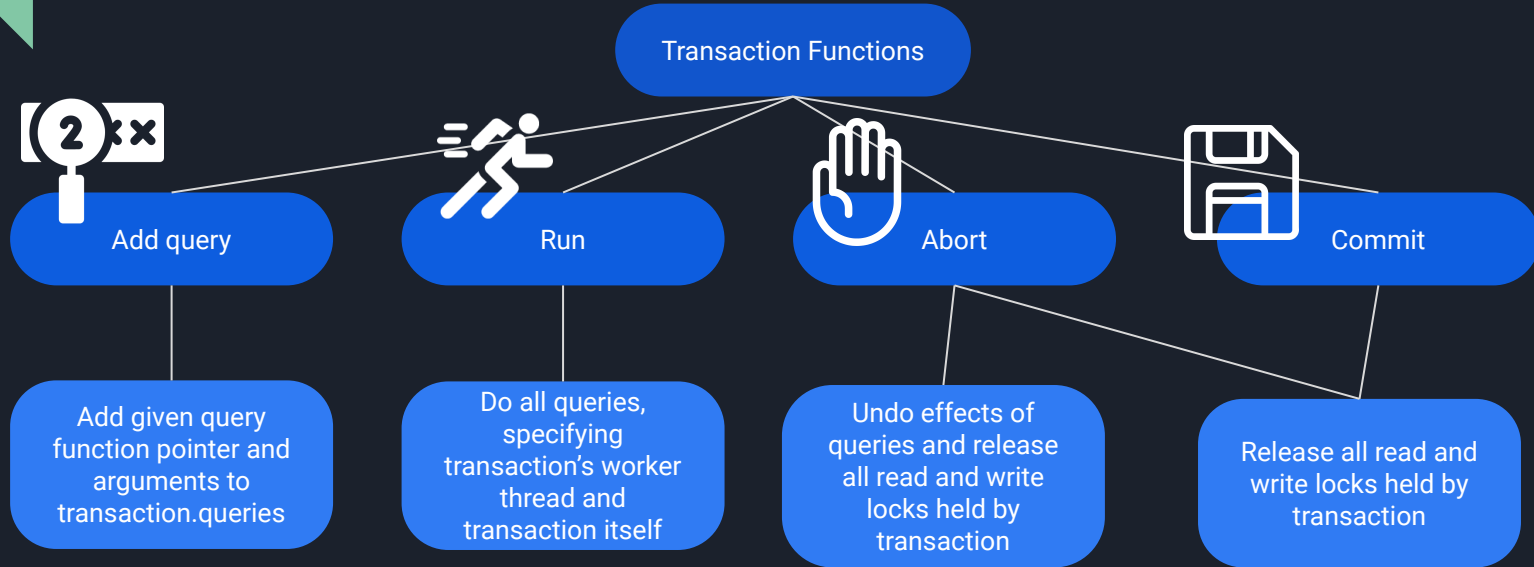
Process Overview (2PL)



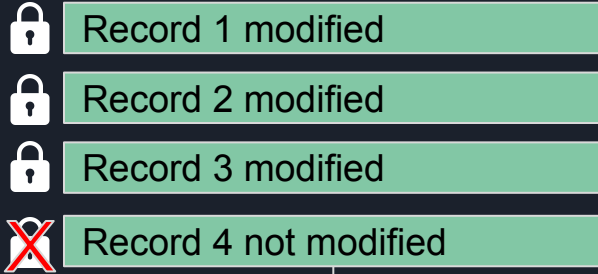
class Transaction (attributes)



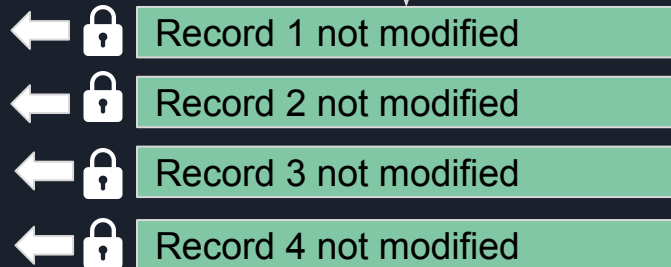
class Transaction (functions)



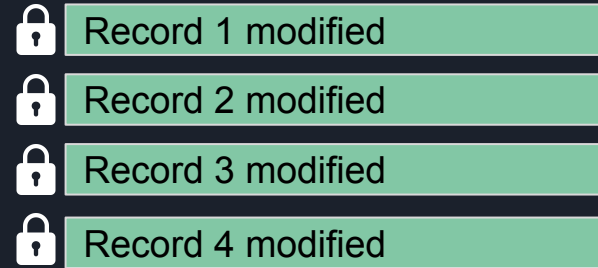
Abort



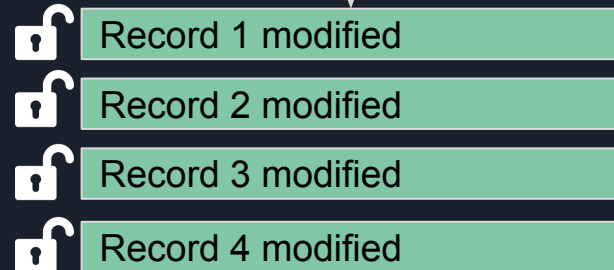
Roll back all updates and
release all locks



Commit

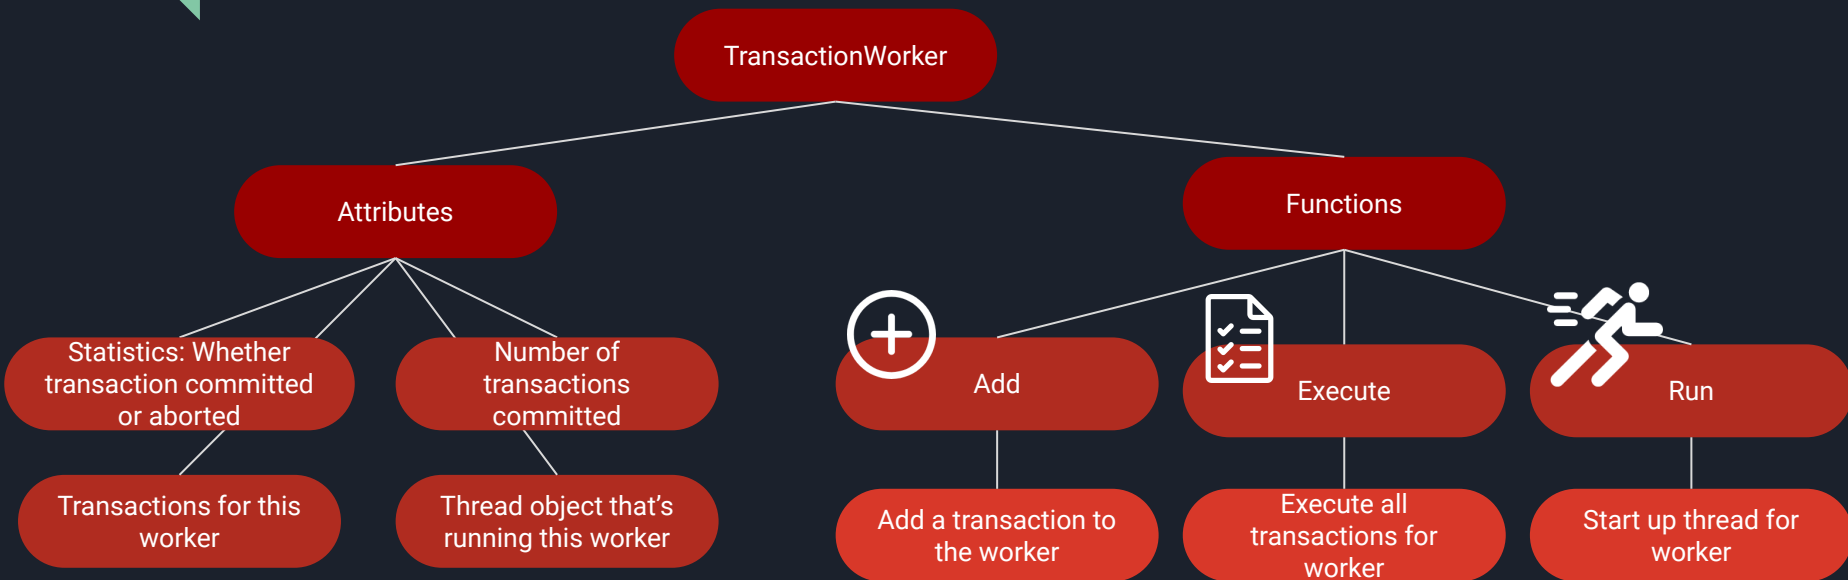


Lock manager releases all
locks on all affected records
for this transaction

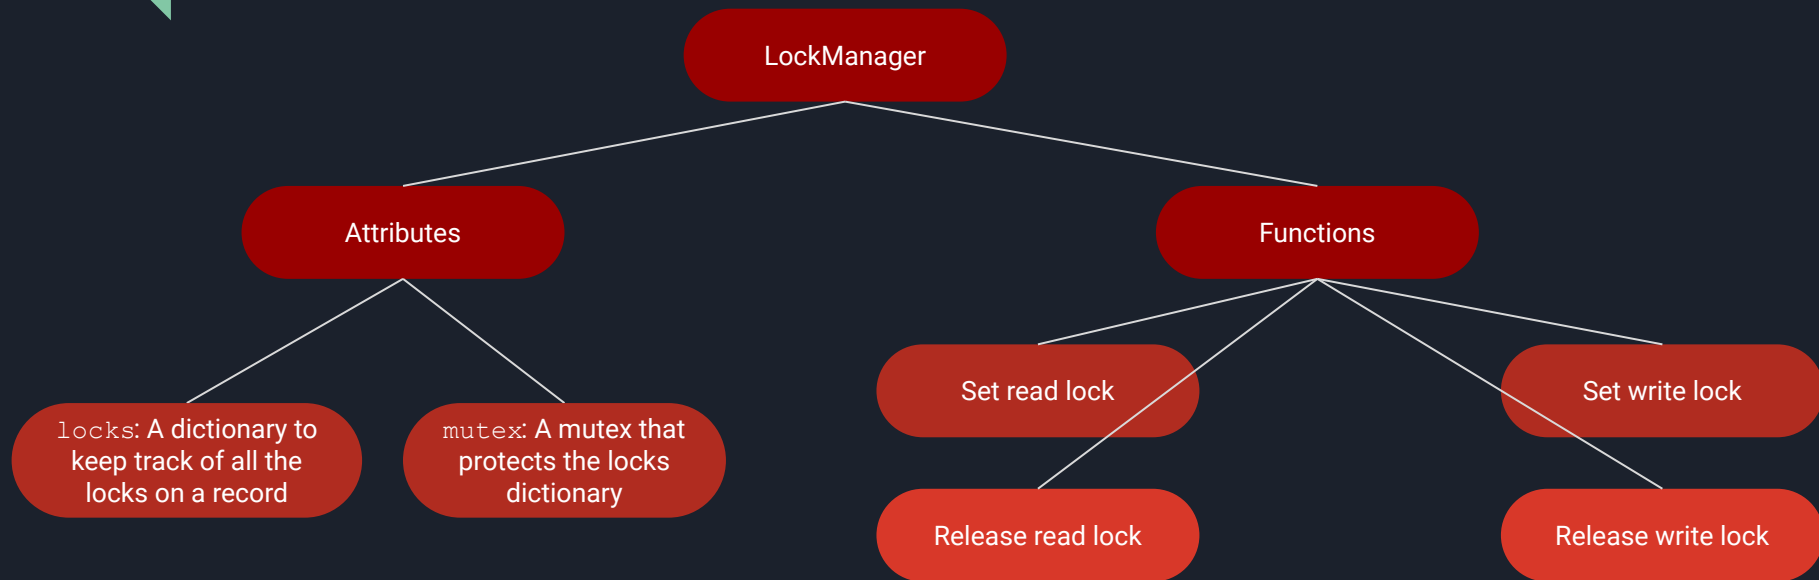




```
class TransactionWorker
```



class LockManager





Lock Management

- `LockManager.locks : {RID: status}`
 - Status: num readers, locked
 - num readers: the number of transactions using an S lock on this RID
 - locked: if the X lock of this RID is taken
- `Transaction.read_locks, Transaction.write_locks (table, RID)`
 - These are stored in a Transaction object because queries in the same transaction shouldn't be restricted by one another's locks
- S locks: select and sum
- X locks: insert, delete, update
- When acquiring a lock for a query, first check `read_locks` and `write_locks` in its transaction. If the RID to be locked isn't found, check with the lock manager for conflicts.
- When a transaction finishes (abort/commit), we remove all the locks used by this transaction from `LockManager.locks`.

Example: acquiring a lock

- `query.update(900, [None, 12, 12, None])`

If RID 900 found,
lock already granted

Update requires an
exclusive lock

Check
`transaction.write_locks`

Check `LockManager.locks`
and `transaction.read_locks`

If locked is True or there are readers
other than the current transaction,
exclusive lock cannot be granted

Set the locked boolean at
`LockManager.locks[900]` to True
Add `(table, 900)` to `transaction.write_locks`



Latches

- `table.id_lock`
 - Protect RID/TID assigning
- `index.locks`
- `Bufferpool.page_locks`
- `LockManager.mutex`
- Since all the read and writes are done in the bf, there's no need to put latches on data files in disk.
- Locks are released when a transaction finishes, but latches are released as soon as the read/write access is finished.
- If a transaction encounters a latch, it will wait till the latch is released to access the protected items instead of aborting.



Merge Update

- Before only merges after closing the database
- Now also merges after a fixed number of updates in a page range
 - The first $n-1$ page ranges can be merged
 - If there is only 1 page range, merge won't happen while database is open
 - Merges all page ranges when closing database however
- 3 Lists to keep track for merge - One data point per page range
 - Page directory list contains the number of merges
 - A new folder created for every merge, numbered sequentially from 0
 - Update list contains the number of updates per page range, updated to zero after merge, initialized at 0
 - TID max list contains the TID cutoff for the indirection column, initialized at new TID

Performance Testing

```
worker_keys = [ {} for t in transaction_workers ]
for i in range(0, 3000):
    key = 92106429 + i
    keys.append(key)
    i = i % num_threads
```

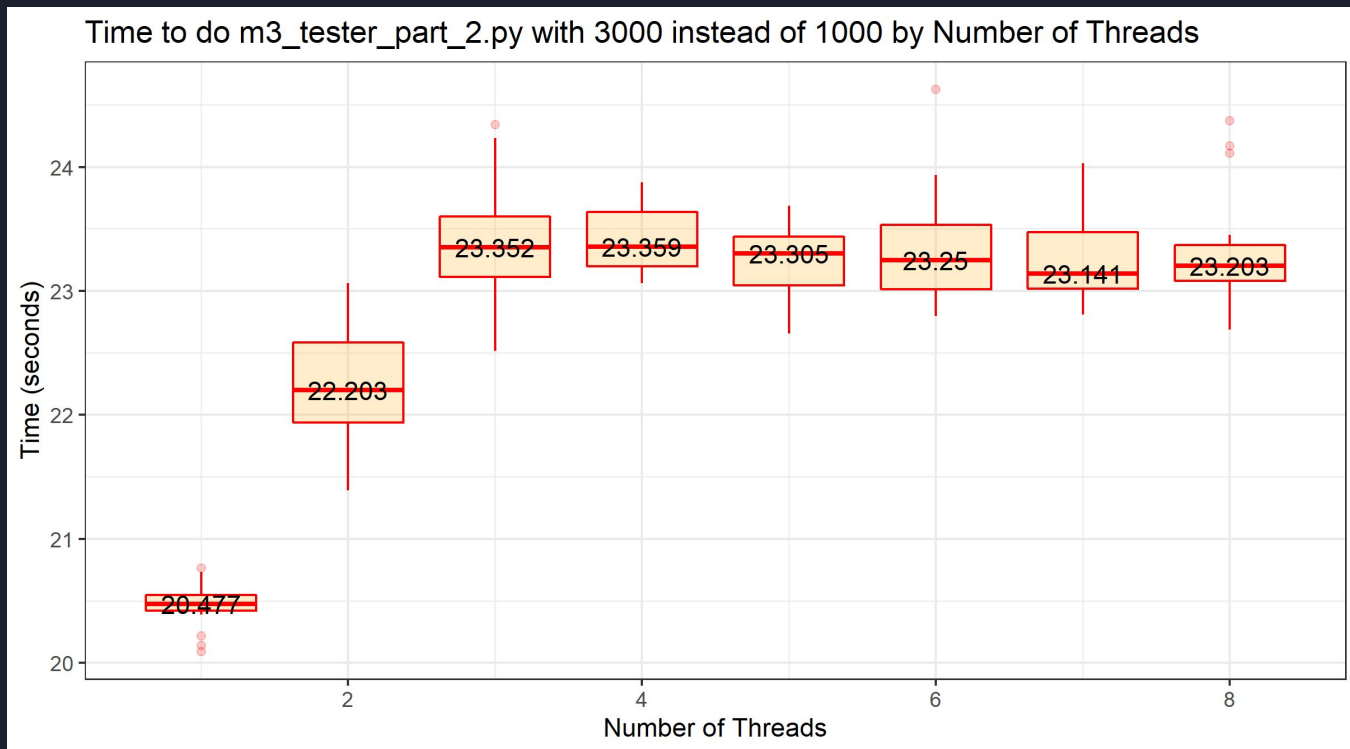
```
start = process_time()
for transaction_worker in transaction_workers:
    transaction_worker.run()

for transaction_worker in transaction_workers:
    transaction_worker.thread.join()

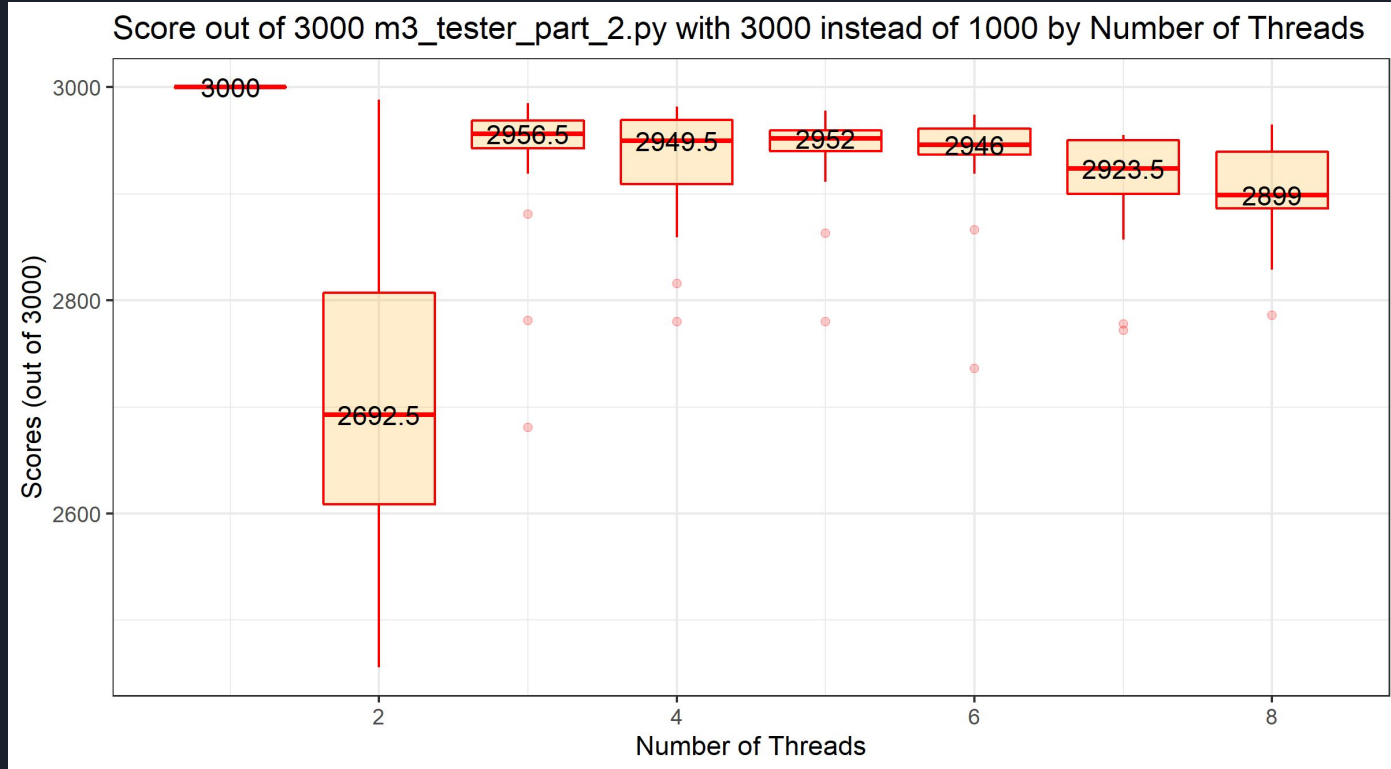
end = process_time()
```

Time to do m3_tester_part_2 by Number of Threads

- Expected to see a concave function, but stabilizes at 23.5 seconds with 3-8 threads
- Side note: The performance testing is done on a slightly different version of the program
 - New version gets higher scores



Score out of 3000 by Number of Threads



Brief Demo: Threads and Abort

